

L10

Metrics

L1 and L2 Regularization

Vanishing / exploding gradients

Activation functions


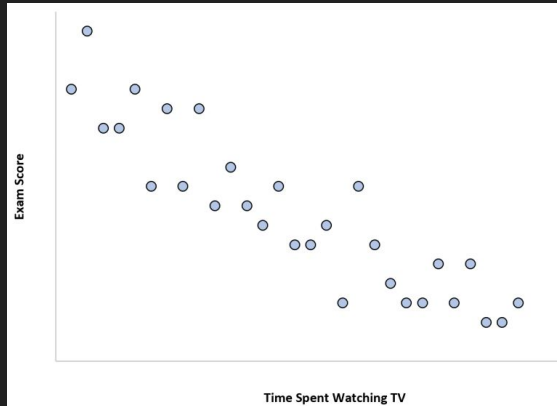
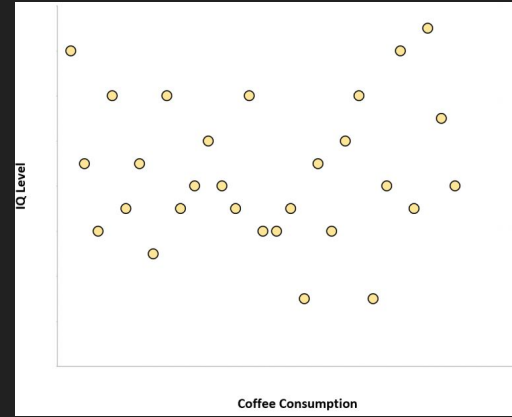

Weight initialization

Metrics https://scikit-learn.org/stable/modules/model_evaluation.html







JAKE-CLARK TUMBLR

Correlation and causation






Correlation does not imply causation




L1 and L2 Regularization

- Lower the weights



Why Lower Weights?



$$\hat{y} = w_0 + w_1x$$

$$\hat{y} = w_0 + w_1x + w_2x^2 + w_3x^3 + w_4x^4$$

$$\hat{y} = w_0 + w_1x + w_2x^2 + w_3x^3 + w_4x^4 + w_5x^5 + w_6x^6 + w_7x^7 + w_8x^8 + w_9x^9 + w_{10}x^{10} + w_{11}x^{11} + w_{12}x^{12} + w_{13}x^{13} + w_{14}x^{14} + w_{15}x^{15}$$

How to Lower Weights?

$$MSE = \sum_i^N (y_i - Wx_i)^2$$

$$MSE_{l1} = \sum_i^N (y_i - Wx_i)^2 + \lambda \sum_j^M |W_j|$$

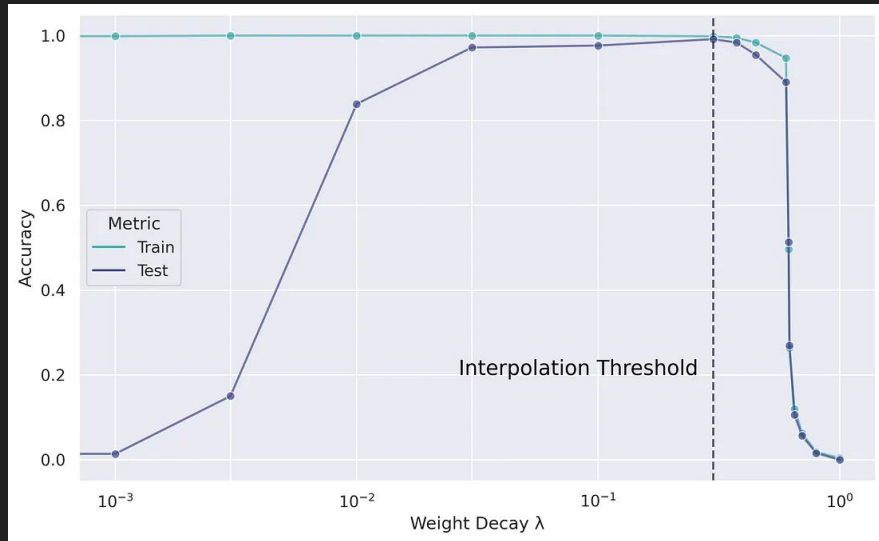

$$MSE_{l2} = \sum_i^N (y_i - Wx_i)^2 + \lambda \sum_j^M W_j^2$$

What if lambda=0, what if lambda=1, what if lambda=inf?

L1 and L2 Regularization


L1 Regularization == L1 Norm == Lasso Regression

L2 Regularization == L2 Norm == Ridge Regression == Weight Decay




Ref: <https://towardsdatascience.com/weight-decay-and-its-peculiar-effects-66e0aee3e7b8>

L1 and L2 changes Loss function topology



L1 vs L2

- L1 tends to generate sparser solutions than a quadratic regularizer




Go Find L2 in Pytorch!

<https://pytorch.org/docs/stable/generated/torch.optim.SGD.html>




Vanishing / exploding gradients




Now, imagine, we have many layers and sigmoid activations ...

$$\frac{\partial l}{\partial w_{1,1}^{(1)}} = \frac{\partial l}{\partial o} \cdot \frac{\partial o}{\partial a_1^{(2)}} \cdot \frac{\partial a_1^{(2)}}{\partial a_1^{(1)}} \cdot \frac{\partial a_1^{(1)}}{\partial w_{1,1}^{(1)}} + \frac{\partial l}{\partial o} \cdot \frac{\partial o}{\partial a_2^{(2)}} \cdot \frac{\partial a_2^{(2)}}{\partial a_1^{(1)}} \cdot \frac{\partial a_1^{(1)}}{\partial w_{1,1}^{(1)}}$$

Vanishing / exploding gradients



\tanh




tanh

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

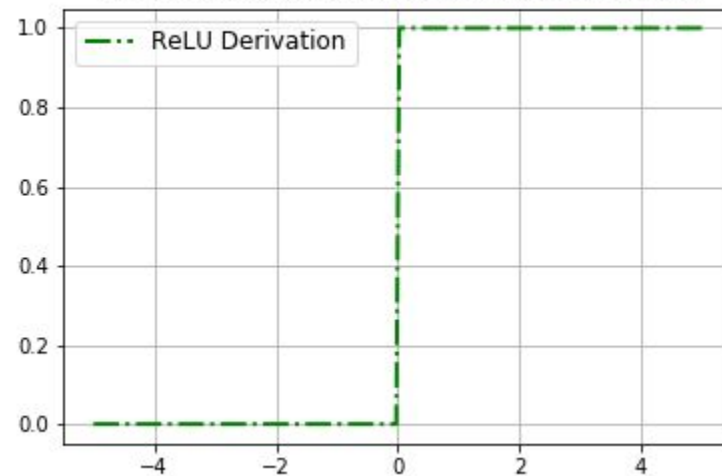
$$\begin{aligned}\frac{d}{dx}\tanh(x) &= \frac{(e^x + e^{-x})(e^x + e^{-x}) - (e^x - e^{-x})(e^x - e^{-x})}{(e^x + e^{-x})^2} \\ &= 1 - \frac{(e^x - e^{-x})^2}{(e^x + e^{-x})^2} = 1 - \tanh^2(x)\end{aligned}$$

ReLU

ReLU Activation function and its derivative




ReLU Activation function and its derivative



ELU

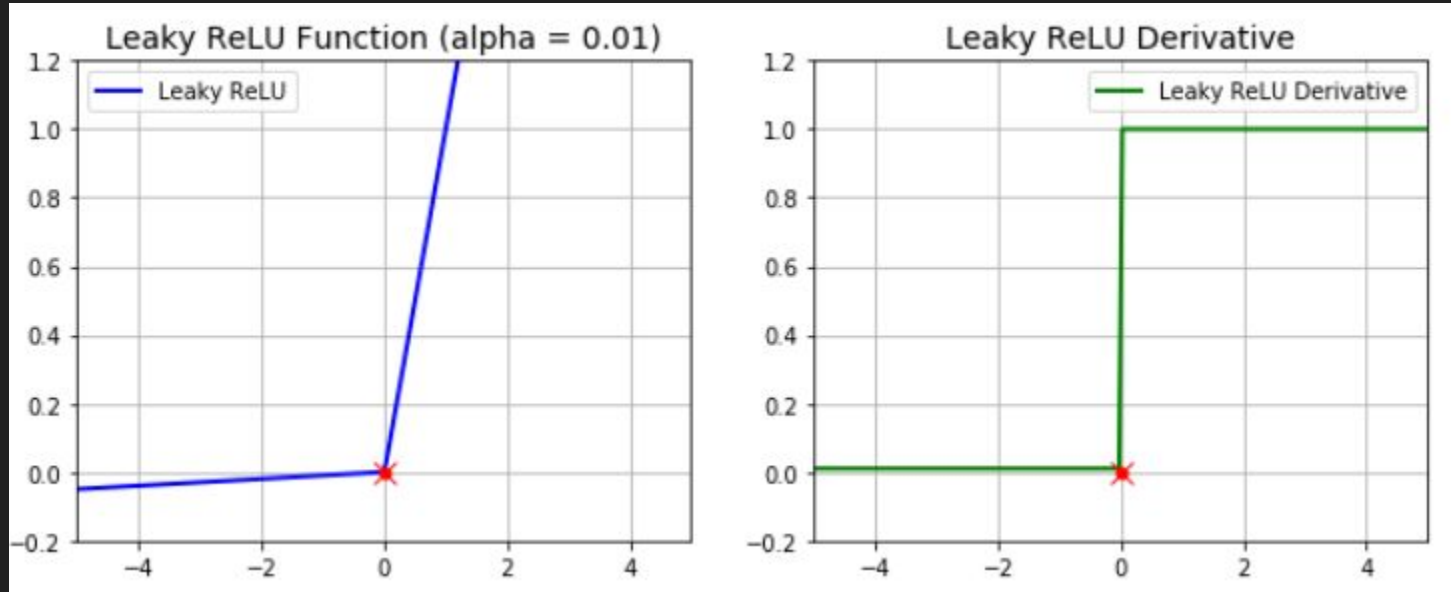
<https://paperswithcode.com/method/elu>

$$f(x) = x \text{ if } x > 0$$
$$\alpha(\exp(x) - 1) \text{ if } x \leq 0$$



Leaky ReLU

<https://paperswithcode.com/method/leaky-relu>




List

Activation Functions


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$

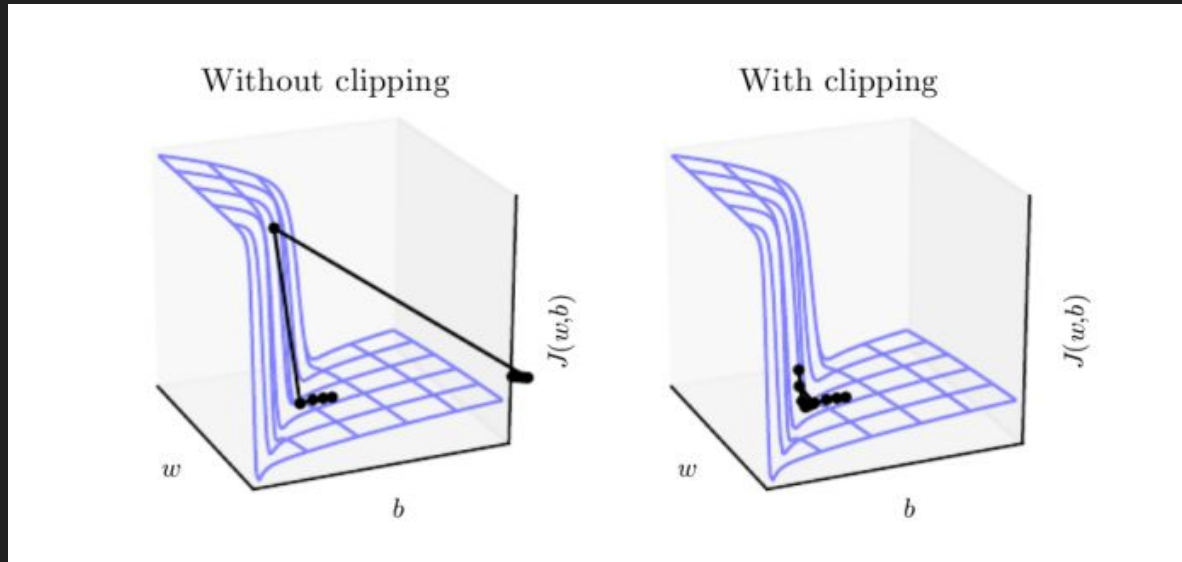


Exploding Gradient. Gradient Clipping

<https://paperswithcode.com/method/gradient-clipping>

https://pytorch.org/docs/stable/generated/torch.nn.utils.clip_grad_norm_.html

<https://github.com/pytorch/pytorch/issues/309#issuecomment-327304962>




Proper Weight Initialization to Solve Vanishing Gradient

- Traditionally, we can initialize weights by sampling from a random uniform distribution in range $[0, 1]$, or better, $[-0.5, 0.5]$
- Or, we could sample from a Gaussian distribution with mean 0 and small variance (e.g., 0.1 or 0.01)

```
 $W^{(l)} = \text{np.random.normal}(loc = 0.0, scale = 1.0) \cdot 0.01$ 
```


Problem with Traditional Weight Initialization



Ref: <https://github.com/ashishpatel26/Tools-to-Design-or-Visualize-Architecture-of-Neural-Network>

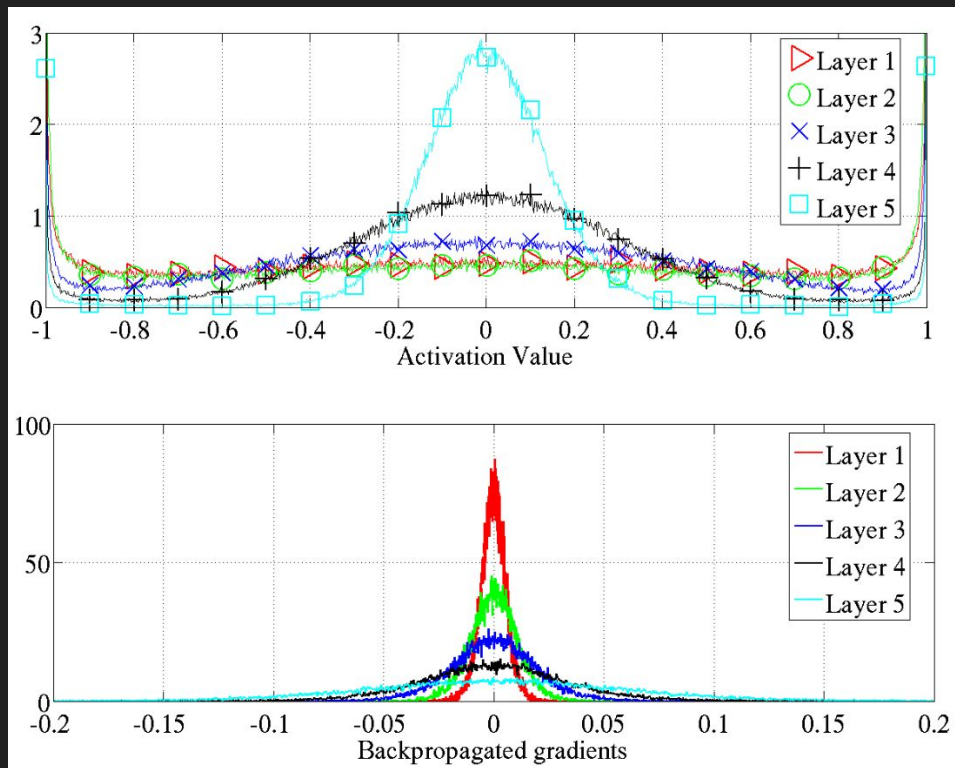
Problem with Traditional Weight Initialization

Gradient Saturation



$$\hat{y} = \sigma(w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + w_5x_5 + w_6x_6 + w_7x_7 + w_8x_8)$$

Vanishing Gradient. Layer-by-Layer



Weight Initialization with Normalization

Xavier Initialization, He Initialization

$$W^{(l)} := W^{(l)} * \sqrt{\frac{gain}{fan_{in}}}$$

$$W^{(l)} := W^{(l)} * \sqrt{\frac{gain}{fan_{out}}}$$

$$W^{(l)} := W^{(l)} * \sqrt{\frac{gain}{fan_{in} + fan_{out}}}$$

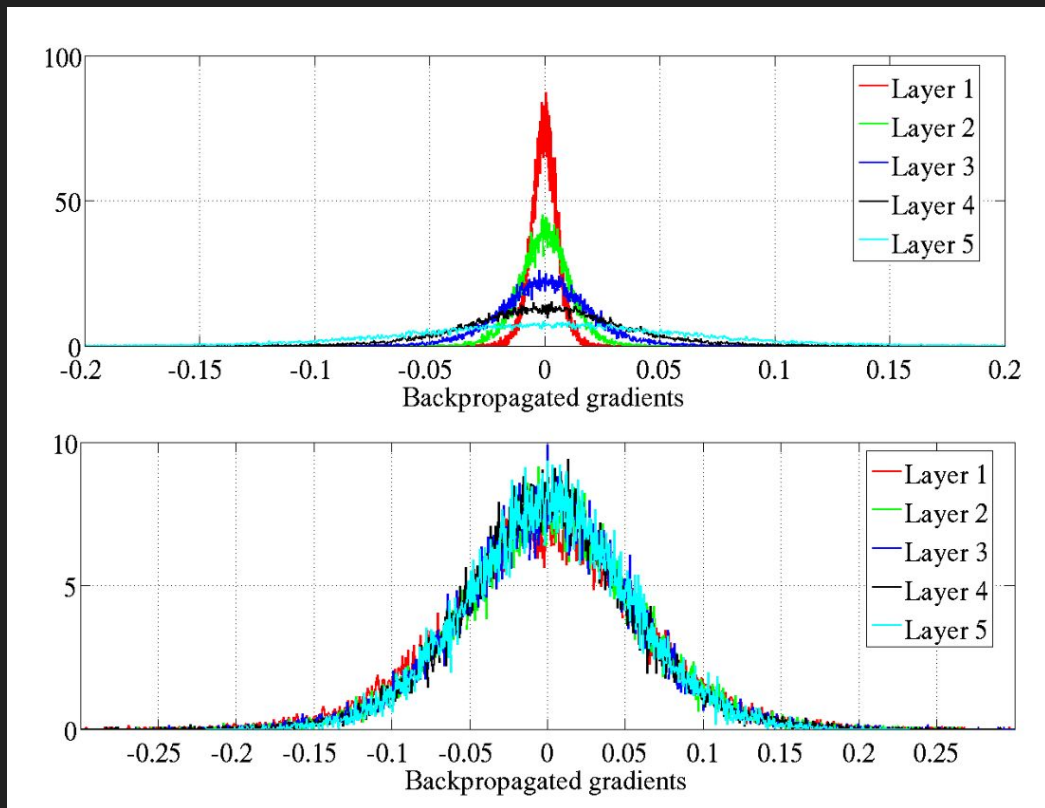
$$W^{(l)} := W^{(l)} * \sqrt{\frac{gain}{(fan_{in} + fan_{out}) / 2}}$$

Weight Initialization

nonlinearity	gain
Linear / Identity	1
Conv{1,2,3}D	1
Sigmoid	1
Tanh	$\frac{5}{3}$
ReLU	$\sqrt{2}$
Leaky Relu	$\sqrt{\frac{2}{1+\text{negative_slope}^2}}$
SELU	$\frac{3}{4}$

Initialization	Activation function	Variance (σ^2)
Glorot	<ul style="list-style-type: none">• Linear• Tanh• Logistic• Softmax	$\sigma^2 = \frac{1}{fan_{avg}}$
He	<ul style="list-style-type: none">• ReLU• Variants of ReLU	$\sigma^2 = \frac{2}{fan_{in}}$
LeCun	<ul style="list-style-type: none">• SELU	$\sigma^2 = \frac{1}{fan_{in}}$

Gradients after Proper Initialization



Note

If BatchNorm is used, initial feature weight choice is less important



Final Notes

Vanishing Gradient:

- Proper weight initialization
- Choose proper activation function
- Architecture tweaks (residual connections)

Gradient Explosion:

- Gradient clipping
- Learning rate normalization (ADAM, LAMB optimizers)



References

- <https://towardsdatascience.com/multi-class-metrics-made-simple-part-ii-the-f1-score-ebe8b2c2ca1>
- <https://neptune.ai/blog/vanishing-and-exploding-gradients-debugging-monitoring-fixing>
- https://ml-cheatsheet.readthedocs.io/en/latest/activation_functions.html
- <https://medium.com/@neuralnets/swish-activation-function-by-google-53e1ea86f820>
- <https://paperswithcode.com/method/swish>

HW

- Apply classification metrics to MNIST classification problem
 - Accuracy (per class and general)
 - Precision (per class and general)
 - Recall (per class and general)
 - F1-score (per class and general)
 - Confusion matrix
 - Classification report