


L12

Computer Vision. Convolutions

MLP for Images

- Can approximate any function (universal approximation theorem)
- Doesn't have local translation invariance, require a lot of data to handle that and a lot of computational resources
- Doesn't care about local neighborhood of a pixel (for example, grass around dog, or sky around tree)




Cat




Cat


Convolution




Convolution



Convolution + Padding



Stride 1 with Padding



Feature Map

Convolution + Padding + Stride

Kernel: 3x3


Stride: 2

Padding: 1

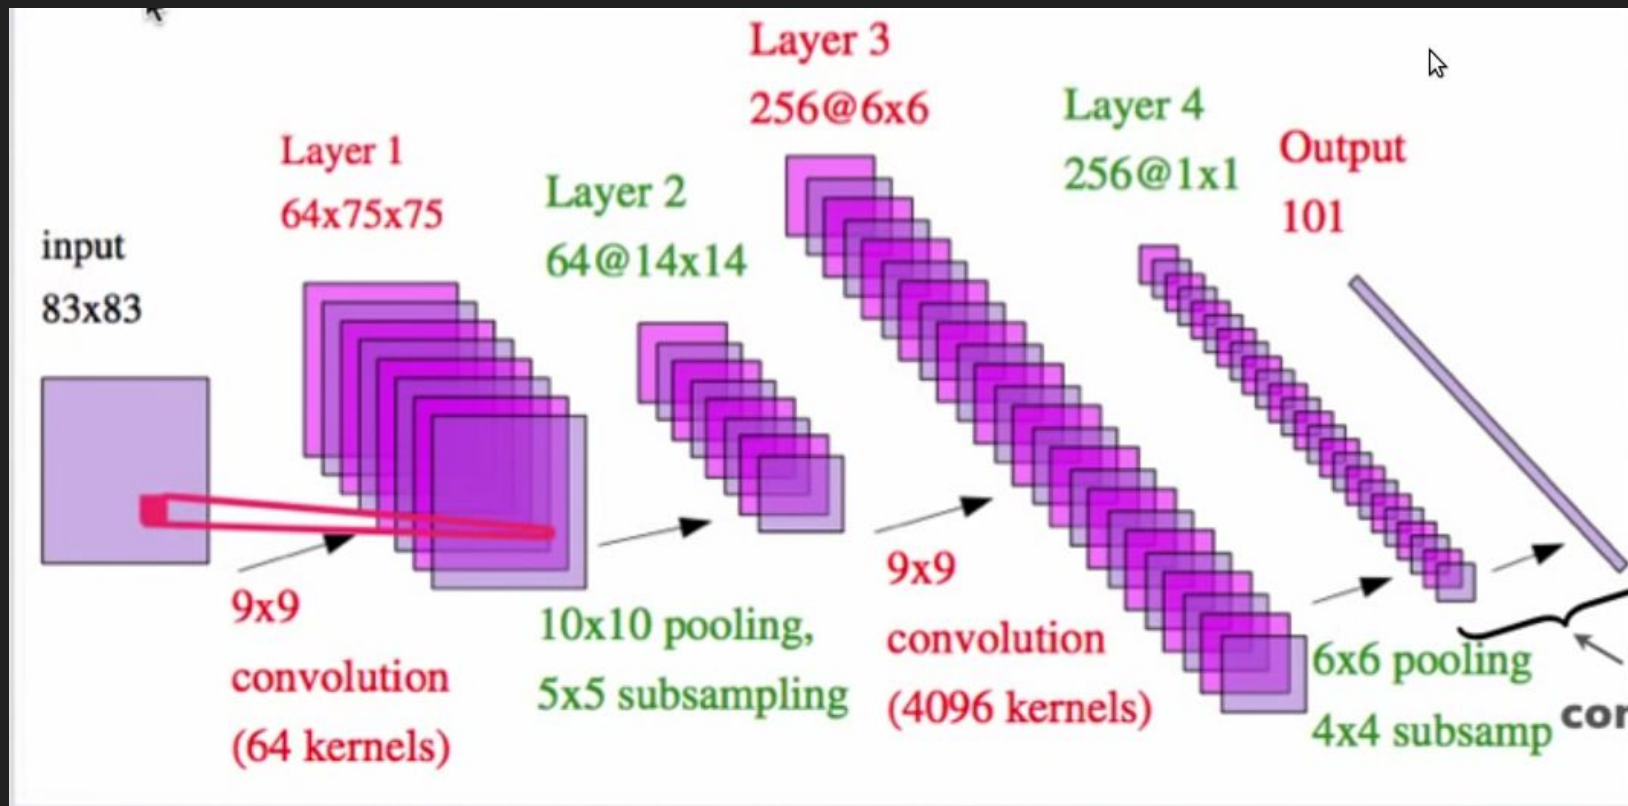
| | | | | | | |
|----------------|----------------|----------------|---|---|---|---|
| 0 ₂ | 0 ₀ | 0 ₁ | 0 | 0 | 0 | 0 |
| 0 ₁ | 2 ₀ | 2 ₀ | 3 | 3 | 3 | 0 |
| 0 ₀ | 0 ₁ | 1 ₁ | 3 | 0 | 3 | 0 |
| 0 | 2 | 3 | 0 | 1 | 3 | 0 |
| 0 | 3 | 3 | 2 | 1 | 2 | 0 |
| 0 | 3 | 3 | 0 | 2 | 3 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 |

| | | |
|---|----|---|
| 1 | 6 | 5 |
| 7 | 10 | 9 |
| 7 | 10 | 8 |

Convolution for each channel


[illegible]

#Filters (kernels)



MaxPool, AvgPool


| | | | |
|---|---|---|---|
| 4 | 3 | 1 | 5 |
| 1 | 3 | 4 | 8 |
| 4 | 5 | 4 | 3 |
| 6 | 5 | 9 | 4 |



| | |
|---|---|
| 4 | 8 |
| 6 | 9 |

Kernel: 2x2
Stride: 1
Padding: 0

| | | | |
|---|---|---|---|
| 4 | 3 | 1 | 5 |
| 1 | 3 | 4 | 8 |
| 4 | 5 | 4 | 3 |
| 6 | 5 | 9 | 4 |



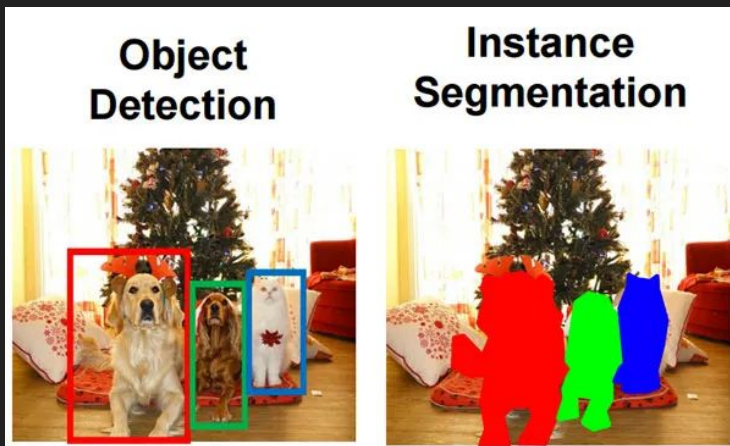
| | |
|-----|-----|
| 2.8 | 4.5 |
| 5.3 | 5.0 |

CV Tasks

- Object detection
- Inpainting
- Segmentation



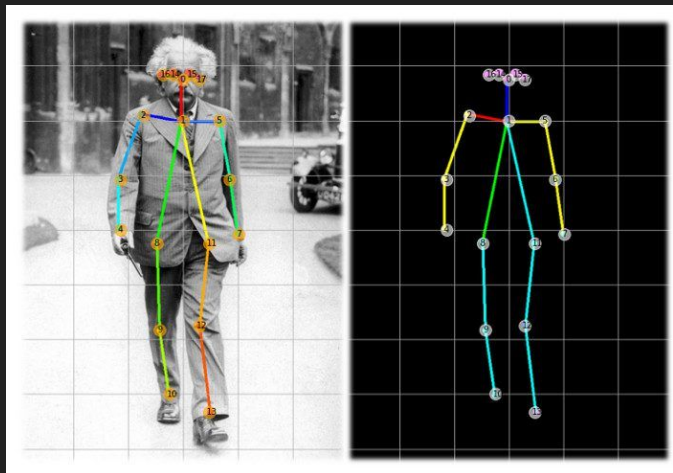


<https://huggingface.co/tasks>



CV Tasks

- Superresolution
- Pose estimation
- Face recognition and authentication



Pre-Convolution Era


Sobel Edge Detection:

$$\begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix}$$

X-Direction Kernel

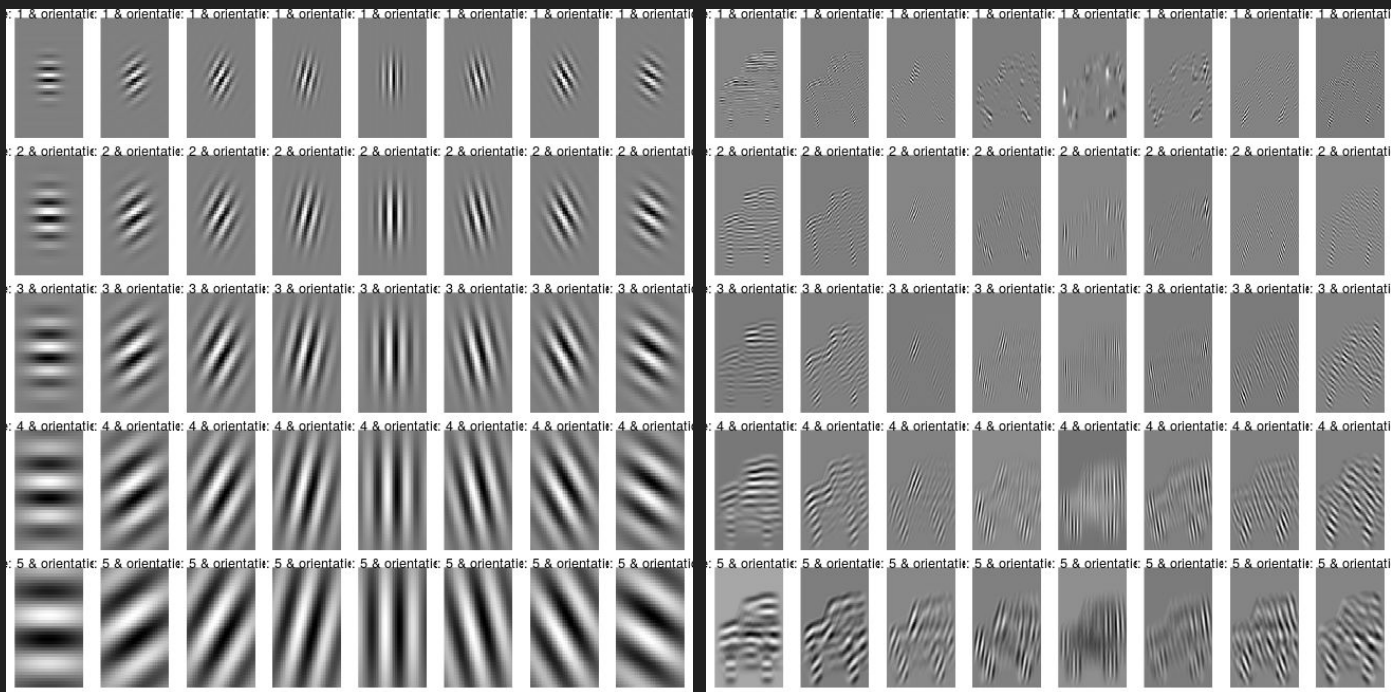
$$\begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Y-Direction Kernel



Gabor Filters


Ref: [link](#), [link](#)




HOG Filters




a Ne



b CS = 3



c CS = 8



d CS = 15




e Su



f CS = 3



g CS = 8



h CS = 15

References

- <https://cs231n.github.io/convolutional-networks/>
- <https://madebyollin.github.io/convnet-calculator/>
- https://scikit-image.org/docs/stable/auto_examples/index.html
- <https://www.kaggle.com/code/isbhargav/guide-to-pytorch-learning-rate-scheduling>
- https://google.github.io/mediapipe/solutions/face_mesh.html
- <https://towardsdatascience.com/hog-histogram-of-oriented-gradients-67ecd887675f>

HW

1. Calculate number of weights on each layer
2. Calculate shape of tensors before and after each layer
3. Make model overfit the data. Show loss curves with overfit
4. Reduce model complexity (number parameters) with keeping accuracy